


## Query Optimization

- Query languages are declarative
- We need to
  - convert declarative statements into executable statements
  - estimate the cost of the executable statements
  - choose the cheapest executable statement among all alternatives

## Query processing stages



## Query Optimization

- In Traditional Databases, we are optimizing for cost of query processing
  - Cost is described in terms of disk access
  - Database keeps statistics about relations, tuples, page sizes.
  - Database also keeps index and sorting information
- Query optimizer
  - uses "cost model" to guess query execution cost for a possible query execution plan
  - chooses a plan which is cheap according to the cost model.

## Relational Data

- Data is Stored in "Tables"
- The size of each "row" in a table is the same.
- Columns can be
  - without index or
  - indexes (B+ tree, Bitmap index, Hash index)
- Columns can be
  - non-clustered
  - clustered (easy access to close/some values)
  - sorted

5

## SQL Query

```
SELECT t1.a, t2.b
  (DISTINCT,SUM(),COUNT())
FROM TABLE1 t1, TABLE2 t2
WHERE t1.a = t2.b AND t1.c > 10
GROUP BY t1.a
```


One Heuristic:

- Apply  $\Pi$ ,  $\sigma$  first
- Apply Joins Next
- Apply Group By Last

```
 $\Pi$  [ $\sigma_{t_1.c > 10}$  (TABLE1  $\bowtie$  TABLE2)]
t1.a
t2.b
 $\Pi$  [ $\sigma_{t_1.a = t_2.b}$  (TABLE1 X TABLE2)]
t1.a
t2.b
 $\Pi$  [ $\sigma_{t_1.c > 10}$  (TABLE1  $\bowtie$  TABLE2)]
t1.a
t2.b
 $\Pi$  [ $\sigma_{t_1.c > 10}$  (TABLE1)]  $\bowtie$   $\Pi$  (TABLE2)]
t1.a a = b t2.b
```


6

Query tree




7



Query tree



8

Query tree





Total cost

Sum of costs of all nodes

If things can be done in parallel, redundant costs can be removed

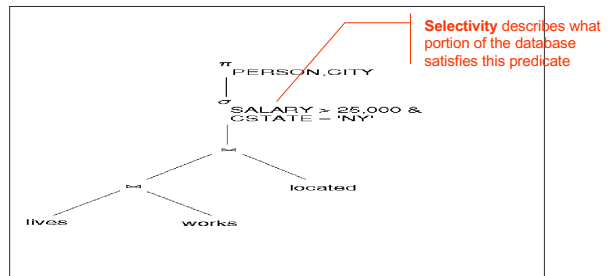
## Query Optimization

- Heuristic-based Query Optimization (CSE 412)
- Cost-based Query Optimization
  - System R Query Optimizer


## Solution:

- Heuristics:
  - perform cheaper operations first
  - cheaper operations eliminate some of the inputs, hence more expensive operations need to deal with a smaller number of inputs.
- Challenge:
  - This assumes that we can freely change the order of operations (associativity, commutativity)

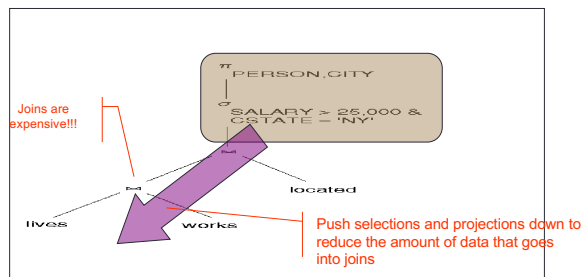
## Query tree



## Query tree



## Query tree



## Cost-based Query Optimization

- Algebraic Transformations
  - commutative
  - distributive, etc.
- Operator Transformation
  - nested loop Vs sort merge join
- Operator insertion

- We are optimizing for cost of query processing
  - Cost is described in terms of disk access
  - Database keeps statistics about relations, tuples, page sizes.
  - Database also keeps index and sorting information
- Query optimizer uses “cost model” to guess query execution cost for a possible query execution plan
  - And chooses a plan which is least costly according to the cost model.

## Cost based query optimization

- Cost Estimation
- Plan enumeration/search space generation
- Search algorithm

## Problems

- What is a good cost model?
- What kind of statistics the database should keep?
- How should the query optimizer create alternative query execution plans?

## Cost estimation

- What is the cost of reading one table from the disc?  
(SCAN)
- What is the cost of doing a selection?  
(FILTER)
- What is the cost of joining two tables?  
(JOIN)

### Statistics:

- We need to estimate the number of disk I/Os.
  - Logical operators / physical operators
  - sizes of intermediate relations
  - ordering of operations (such as joins)
- How do we estimate cost of operations and sizes of resulting relations?
  - $B(R) \leftarrow$  number of blocks
  - $T(R) \leftarrow$  number of tuples
  - $V(R.a) \leftarrow$  number of distinct values for an attribute "a"
  - $V(R.[a_1 \dots a_n]) \leftarrow$  number of distinct values when  $a_1 \dots a_n$  are considered together.

- Cost Estimates
  - I/O cost + CPU cost
- Scan Cost
  - $B(R)$  (unsorted)
  - $B(R) + \# \text{ of index pages}$  (sorted on a clustered index)
  - $T(R) + \# \text{ of index pages}$  (sorted on a non-clustered index)
- (We already talked about the affects of database buffers)

### Example use of statistics:

- Projection
  - easier
    - omit the attributes that are not included in the projection
    - recalculate how many blocks would be needed to store the resulting tuples.
  - What if distinct is specified?
- Selection
  - $S = \sigma_{a=c}(R)$  when "c" is a constant
    - $T(S) = \frac{T(R)}{V(R.a)}$
    - $S = \sigma_{a < c}(R)$  when "c" is a constant
      - $T(S) = \frac{T(R)}{V(R.a)} [V(R.a) - 1]$  or  $T(S) = T(R) \quad \text{????}$

- $S = \sigma_{c_1 \text{ or } c_2}(R)$ 
  - Assume  $c_1$  and  $c_2$  are independent
  - $T(S) = T(R) * [1 - (1 - \frac{n}{T(R)}) * (1 - \frac{m}{T(R)})]$ ,
    - where "n" is the number of tuples satisfying  $c_1$ , and "m" is the number of tuples satisfying  $c_2$

- Join

- $R(x,y) \bowtie S(y,z)$ 
  - $T(R \bowtie S) = 0$
  - $T(R \bowtie S) = T(R)$
  - $T(R \bowtie S) = T(R)T(S)$
- $T(R \bowtie S) = \frac{T(R)T(S)}{\max(V(R,y), V(S,y))}$
- $R(x,y_1, y_2) \bowtie S(y_1, y_2, z)$
- $T(R \bowtie S) = \frac{T(R)T(S)}{\max(V(R,y_1)V(S, y_1))\max(V(R, y_2), V(S, y_2))}$

- Alternative: Keep Histograms!

- Join-cost

- Nested Loop join  
 $\text{scan}(R_1) + T(R_1) * \text{scan}(R_2)$
- Sort-merge join  
 $\text{sorted-scan}(R_1) + \text{sorted-scan}(R_2)$

## Assumptions of cost-based query optimization

- Independent of how a query is executed we get the same number of results:

$$\sigma_{\theta_1 \wedge \theta_2}(R) = \sigma_{\theta_1}(\sigma_{\theta_2}(R)) = \sigma_{\theta_2}(\sigma_{\theta_1}(R))$$

## Assumptions of cost-based query optimization

- Principle of sub-query optimality: An optimal plan for a query includes optimal subplans for the subqueries

## Assumptions of cost-based query optimization

- Principle of sub-query optimality: An optimal plan for a query includes optimal subplans for the subqueries

$$\begin{aligned}\text{cost}^1(\sigma_{\theta_1}(\sigma^1_{\theta_2}(R))) &= \text{cost}(\sigma^1_{\theta_2}(R)) + f(\text{size}(\sigma^1_{\theta_2}(R)), \theta_1) \\ \text{cost}^2(\sigma_{\theta_1}(\sigma^2_{\theta_2}(R))) &= \text{cost}(\sigma^2_{\theta_2}(R)) + f(\text{size}(\sigma^2_{\theta_2}(R)), \theta_1)\end{aligned}$$

## Assumptions of cost-based query optimization

- Principle of sub-query optimality: An optimal plan for a query includes optimal subplans for the subqueries

$$\begin{aligned}\text{cost}^1(\sigma_{\theta_1}(\sigma^1_{\theta_2}(R))) &= \text{cost}(\sigma^1_{\theta_2}(R)) + f(\text{size}(\sigma^1_{\theta_2}(R)), \theta_1) \\ &= \\ \text{cost}^2(\sigma_{\theta_1}(\sigma^2_{\theta_2}(R))) &= \text{cost}(\sigma^2_{\theta_2}(R)) + f(\text{size}(\sigma^2_{\theta_2}(R)), \theta_1)\end{aligned}$$

## Assumptions of cost-based query optimization

- Principle of sub-query optimality: An optimal plan for a query includes optimal subplans for the subqueries

$$\begin{aligned}\text{cost}^1(\sigma_{\theta_1}(\sigma^1_{\theta_2}(R))) &= \text{cost}(\sigma^1_{\theta_2}(R)) + f(\text{size}(\sigma^1_{\theta_2}(R)), \theta_1) \\ &= \\ \text{cost}^2(\sigma_{\theta_1}(\sigma^2_{\theta_2}(R))) &= \text{cost}(\sigma^2_{\theta_2}(R)) + f(\text{size}(\sigma^2_{\theta_2}(R)), \theta_1)\end{aligned}$$

## Assumptions of cost-based query optimization

- Principle of sub-query optimality: An optimal plan for a query includes optimal subplans for the subqueries

$$\begin{aligned}\text{cost}^1(\sigma_{\theta_1}(\sigma^1_{\theta_2}(R))) &= \text{cost}(\sigma^1_{\theta_2}(R)) + f(\text{size}(\sigma^1_{\theta_2}(R)), \theta_1) \\ &= \\ \text{cost}^2(\sigma_{\theta_1}(\sigma^2_{\theta_2}(R))) &= \text{cost}(\sigma^2_{\theta_2}(R)) + f(\text{size}(\sigma^2_{\theta_2}(R)), \theta_1)\end{aligned}$$

Pick the cheapest!!!



## Assumptions of cost-based query optimization


- Principle of sub-query optimality: An optimal plan for a query includes optimal subplans for the subqueries
- ...we can use recursion!!!!!!!!!!

### Example:


$R_1(a,b,c) \bowtie R_2(a,d) \bowtie R_3(d,e) \bowtie R_4(c,e)$

- Question: Which order of join is the best?
- Important:
  - The join ordering problem is NP-Hard!
    - I.e., there is no known polynomial time algorithm
    - There is strong reasons to believe that there is no polynomial time algorithm

## What about join orders???



## What about join orders???



The goal of query optimization is to **eliminate costliest** executions....  
 .....finding cheapest executions is very very expensive.

## Recursion

- Given  $Q = R_1 \times R_2 \times R_3 \times \dots \times R_n$ 
  - for all  $R_i$ ,
    - find the cheapest plan for  $Q_{(-R_i)}$
    - compute  $\text{cost}_i = \text{cost}(Q_{(-R_i)} \times R_i)$
  - Pick the plan with the smallest cost


## Recursion

- Given  $Q = R_1 \times R_2 \times R_3 \times \dots \times R_n$ 
    - for all  $R_i$ ,
      - find the cheapest plan for  $Q_{(-R_i)}$
      - compute  $\text{cost}_i = \text{cost}(Q_{(-R_i)} \times R_i)$
    - Pick the plan with the smallest cost
- Recursion!!!

## Recursion

- Given  $Q = R_1 \times R_2 \times R_3 \times \dots \times R_n$ 
  - for all  $R_i$ ,
    - find the cheapest plan for  $Q_{(-R_i)}$
    - compute  $\text{cost}_i = \text{cost}(Q_{(-R_i)} \times R_i)$
  - Pick the plan with the smallest cost
- There is a dynamic programming based algorithm that uses this recursion.

## Dynamic Programming



## Additional heuristics

- Avoid cartesian products
  - No index, hash, or sort-merge opportunity
  - The output is large
- Do not maintain one single plan per subexpression, but maintain multiple plans for “interesting” orders
  - If the result of the subexpression is sorted on the attributes specified in a sort operator in the query
  - If the attributes are used in a group\_by operator
  - If they are used in a later join

## Rule-based optimizers

- Starburst
  - Query is represented as a graph
  - If-C-then-A rules are used to apply algebraic transformations
  - A separate plan optimizer is used for physical query optimization

## Rule-based optimizers

- Volcano
  - Query is represented as a tree
  - All transformations (logical, physical) treated uniformly
  - It uses a top-down query optimization algorithm (branch and bound)
    - Subexpressions are optimized on demand

## Branch-and-bound

- Use heuristics to find a good physical plan
  - (let its cost be  $C$ )
- Refine the plan as follows
  - Consider other plans for subqueries
  - Any subquery with cost  $>C$  is ignored
  - If you find a complete plan with cost  $C' < C$  then replace the plan.

## Greedy Algorithm

- Make a decision at a time about joins
  - Never backtrack

$$s = \frac{\text{size of the join result}}{\text{size of the current result}}$$

- Always pick the option with smallest s