



# Crash Recovery

## Chapter 18




# Review: The ACID properties

- ❖ **A**tomicity: All actions in the Xact happen, or none happen.
- ❖ **C**onsistency: If each Xact is consistent, and the DB starts consistent, it ends up consistent.
- ❖ **I**solation: Execution of one Xact is isolated from that of other Xacts.
- ❖ **D**urability: If a Xact commits, its effects persist.
- ❖ The **Recovery Manager** guarantees Atomicity & Durability.



# Motivation


- ❖ Atomicity:
  - Transactions may abort ("Rollback").
- ❖ Durability:
  - What if DBMS stops running? (Causes?)
- ❖ Desired Behavior after system restarts:
  - T1, T2 & T3 should be durable.
  - T4 & T5 should be aborted (effects not seen).



# Assumptions

- ❖ Concurrency control is in effect.
  - **Strict 2PL**, in particular.
- ❖ Updates are happening "in place".
  - i.e. data is overwritten on (deleted from) the disk.
- ❖ A simple scheme to guarantee Atomicity & Durability?

## Handling the Buffer Pool



### ❖ **Force** every write to disk?

- Poor response time.
- But provides durability.

### ❖ **Steal** buffer-pool frames from uncommitted Xacts?

- If not, poor throughput.
- If so, how can we ensure atomicity?

	No Steal	Steal
Force	Trivial	
No Force		Desired

## More on Steal and Force




### ❖ **STEAL** (why enforcing Atomicity is hard)

- *To steal frame F:* Current page in F (say P) is written to disk; some Xact holds lock on P.
  - What if the Xact with the lock on P aborts?
  - Must remember the old value of P at steal time (to support **UNDO**ing the write to page P).

### ❖ **NO FORCE** (why enforcing Durability is hard)

- What if system crashes before a modified page is written to disk?
- Write as little as possible, in a convenient place, at commit time, to support **REDO**ing modifications.

## Basic Idea: Logging



### ❖ Record REDO and UNDO information, for every update, in a **log**.

- Sequential writes to log (put it on a separate disk).
- Minimal info (diff) written to log, so multiple updates fit in a single log page.

### ❖ **Log:** An ordered list of REDO/UNDO actions

- Log record contains:
  - <XID, pageID, offset, length, old data, new data>
- and additional control info (which we'll see soon).

## Write-Ahead Logging (WAL)



### ❖ The **Write-Ahead Logging** Protocol:

- ① Must **force** the **log record** for an update **before** the corresponding **data page** gets to disk.
- ② Must **write all log records** for a Xact **before commit**.


### ❖ #1 guarantees Atomicity.

### ❖ #2 guarantees Durability.

### ❖ Exactly how is logging (and recovery!) done?

- We'll study the ARIES algorithms.

## WAL & the Log



- ❖ Each log record has a unique **Log Sequence Number (LSN)**.

- LSNs always increasing.

- ❖ Each **data page** contains a **pageLSN**.


- The LSN of the most recent *log record* for an update to that page.

- ❖ System keeps track of **flushedLSN**.

- The max LSN flushed so far.

- ❖ **WAL: Before** a page is written,

- $\text{pageLSN} \leq \text{flushedLSN}$



## Log Records



### LogRecord fields:

prevLSN  
 XID  
 type  
 pageID  
 length  
 offset  
 before-image  
 after-image

update records only

### Possible log record types:

- ❖ **Update**
- ❖ **Commit**
- ❖ **Abort**
- ❖ **End** (signifies end of commit or abort)
- ❖ **Compensation Log Records (CLRs)**
  - for UNDO actions

## Other Log-Related State



### ❖ Transaction Table:

- One entry per active Xact.
- Contains **XID**, **status** (running/committed/aborted), and **lastLSN**.

### ❖ Dirty Page Table:

- One entry per dirty page in buffer pool.
- Contains **recLSN** -- the LSN of the log record which **first** caused the page to be dirty.

## Normal Execution of an Xact



- ❖ Series of **reads & writes**, followed by **commit** or **abort**.

- We will assume that write is atomic on disk.
  - In practice, additional details to deal with non-atomic writes.

### ❖ Strict 2PL.

- ❖ STEAL, NO-FORCE buffer management, with **Write-Ahead Logging**.

## Checkpointing



- ❖ Periodically, the DBMS creates a checkpoint, in order to minimize the time taken to recover in the event of a system crash. Write to log:
  - **begin\_checkpoint** record: Indicates when chkpt began.
  - **end\_checkpoint** record: Contains current *Xact table* and *dirty page table*. This is a 'fuzzy checkpoint':
    - Other Xacts continue to run; so these tables accurate only as of the time of the **begin\_checkpoint** record.
    - No attempt to force dirty pages to disk; effectiveness of checkpoint limited by oldest unwritten change to a dirty page. (So it's a good idea to periodically flush dirty pages to disk!)
  - Store LSN of chkpt record in a safe place (*master* record).

## The Big Picture: What's Stored Where



### LogRecords

prevLSN  
XID  
type  
pageID  
length  
offset  
before-image  
after-image



### Data pages

each  
with a  
pageLSN

### master record



### Xact Table

lastLSN  
status

### Dirty Page Table

recLSN


### flushedLSN

## Simple Transaction Abort



- ❖ For now, consider an explicit abort of a Xact.
  - No crash involved.
- ❖ We want to "play back" the log in reverse order, UNDOing updates.
  - Get **lastLSN** of Xact from Xact table.
  - Can follow chain of log records backward via the **prevLSN** field.
  - Before starting UNDO, write an **Abort log record**.
    - For recovering from crash during UNDO!

## Abort, cont.




- ❖ To perform UNDO, must have a lock on data!
  - No problem!
- ❖ Before restoring old value of a page, write a CLR:
  - You continue logging while you UNDO!!
  - CLR has one extra field: **undonextLSN**
    - Points to the next LSN to undo (i.e. the prevLSN of the record we're currently undoing).
  - CLRs *never* Undone (but they might be Redone when repeating history: guarantees Atomicity!)
- ❖ At end of UNDO, write an "end" log record.

## Transaction Commit



- ❖ Write **commit** record to log.
- ❖ All log records up to Xact's **lastLSN** are flushed.
  - Guarantees that **flushedLSN**  $\geq$  **lastLSN**.
  - Note that log flushes are sequential, synchronous writes to disk.
  - Many log records per log page.
- ❖ Commit() returns.
- ❖ Write **end** record to log.

## Crash Recovery: Big Picture



- ❖ Start from a **checkpoint** (found via **master** record).
- ❖ Three phases. Need to:
  - Figure out which Xacts committed since checkpoint, which failed (**Analysis**).
  - **REDO** *all* actions.
    - ♦ (repeat history)
  - **UNDO** effects of failed Xacts.

## Recovery: The Analysis Phase



- ❖ Reconstruct state at checkpoint.
  - via **end\_checkpoint** record.
- ❖ Scan log forward from checkpoint.
  - **End** record: Remove Xact from Xact table.
  - **Other records**: Add Xact to Xact table, set **lastLSN=LSN**, change Xact status on **commit**.
  - **Update** record: If P not in Dirty Page Table,
    - Add P to D.P.T., set its **recLSN=LSN**.

## Recovery: The REDO Phase



- ❖ We **repeat History** to reconstruct state at crash:
  - Reapply *all* updates (even of aborted Xacts!), redo CLR's.
- ❖ Scan forward from log rec containing smallest **recLSN** in D.P.T. For each CLR or update log rec **LSN**, REDO the action unless:
  - Affected page is not in the Dirty Page Table, or
  - Affected page is in D.P.T., but has **recLSN**  $>$  **LSN**, or
  - **pageLSN** (in DB)  $\geq$  **LSN**.
- ❖ To **REDO** an action:
  - Reapply logged action.
  - Set **pageLSN** to **LSN**. No additional logging!

## Recovery: The UNDO Phase

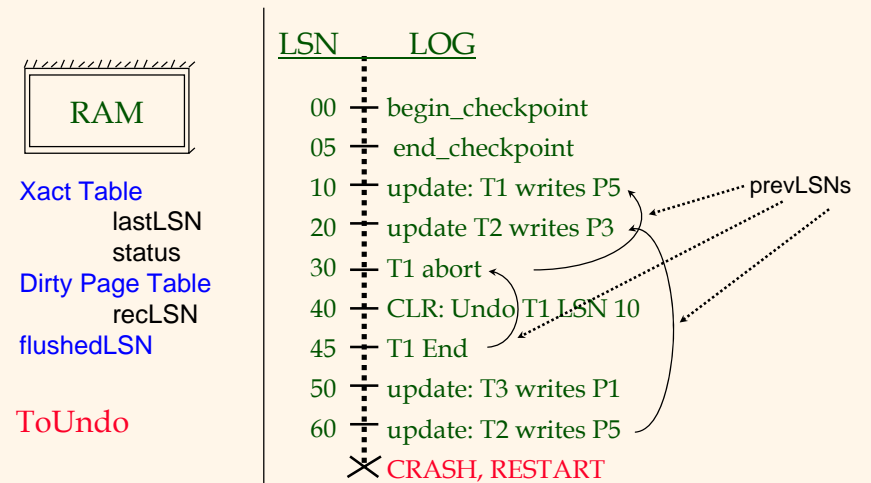
$ToUndo = \{ l \mid l \text{ a lastLSN of a "loser" Xact} \}$

### Repeat:

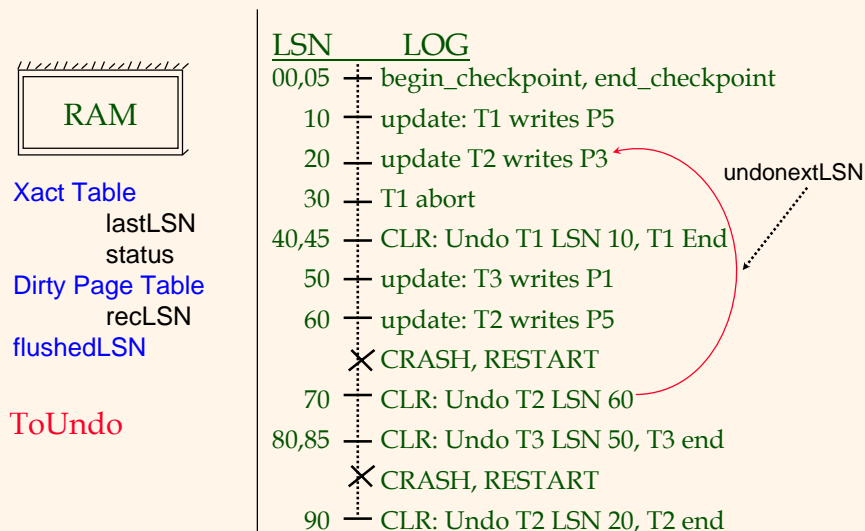
- Choose largest LSN among ToUndo.
- If this LSN is a **CLR** and **undonextLSN == NULL**
  - Write an **End** record for this Xact.
- If this LSN is a **CLR**, and **undonextLSN != NULL**
  - Add **undonextLSN** to ToUndo
- Else this LSN is an **update**. Undo the update, write a CLR, add **prevLSNs** to ToUndo.

Until **ToUndo** is empty.

## Example of Recovery



## Example: Crash During Restart!



## Additional Crash Issues

- ❖ What happens if system crashes during Analysis? During REDO?
- ❖ How do you limit the amount of work in REDO?
  - Flush asynchronously in the background.
  - Watch "hot spots"!
- ❖ How do you limit the amount of work in UNDO?
  - Avoid long-running Xacts.

## Summary of Logging/Recovery



- ❖ **Recovery Manager** guarantees Atomicity & Durability.
- ❖ Use WAL to allow STEAL/NO-FORCE w/o sacrificing correctness.
- ❖ LSNs identify log records; linked into backwards chains per transaction (via prevLSN).
- ❖ pageLSN allows comparison of data page and log records.

## Summary, Cont.



- ❖ **Checkpointing**: A quick way to limit the amount of log to scan on recovery.
- ❖ Recovery works in 3 phases:
  - **Analysis**: Forward from checkpoint.
  - **Redo**: Forward from oldest recLSN.
  - **Undo**: Backward from end to first LSN of oldest Xact alive at crash.
- ❖ Upon Undo, write CLR.
- ❖ Redo “repeats history”: Simplifies the logic!